

# Project Report: Task 1 - Titanic: Machine Learning from Disaster

## 1. Introduction

### 1.1 Project Overview

The objective of this project is to build a predictive model that classifies whether a passenger survived the Titanic disaster based on their features such as age, gender, class, and other relevant attributes. This problem is a classic example of a binary classification task in machine learning.

### 1.2 Dataset Description

The dataset used in this project is the Titanic dataset, which includes various features for each passenger:

- **PassengerId:** Unique identifier for each passenger.
- **Survived:** Survival status (0 = No, 1 = Yes).
- **Pclass:** Ticket class (1 = 1st, 2 = 2nd, 3 = 3rd).
- **Name:** Passenger's name.
- **Sex:** Passenger's sex.
- **Age:** Passenger's age.
- **SibSp:** Number of siblings/spouses aboard the Titanic.
- **Parch:** Number of parents/children aboard the Titanic.
- **Ticket:** Ticket number.
- **Fare:** Passenger fare.
- **Cabin:** Cabin number.
- **Embarked:** Port of embarkation (C = Cherbourg, Q = Queenstown, S = Southampton).

## 2. Data Exploration

### 2.1 Understanding the Dataset

Explored the dataset to understand its structure, feature distribution, and identify any potential issues such as missing values or outliers.

### 2.2 Checking for Missing Data and Outliers

- **Missing Data:**
  - Checked for missing values in each feature.
  - Found missing values in 'Age', 'Cabin', and 'Embarked' columns.
- **Outliers:**
  - Visualized the distribution of numerical features (e.g., 'Age', 'Fare') to identify outliers.

```
python
Copy code
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

# Load the datasets
titanic_data = pd.read_csv('Task 1 Titanic Machine Learning from Disaster/train.csv')

# Checking for missing values
missing_values = titanic_data.isnull().sum()
print(missing_values)

# Visualizing outliers in Age and Fare
sns.boxplot(x=titanic_data['Age'])
plt.show()

sns.boxplot(x=titanic_data['Fare'])
plt.show()
```

## 3. Data Preprocessing

### 3.1 Handling Missing Values

- **Age:** Imputed missing values using the median age.
- **Cabin:** Dropped this feature due to a high percentage of missing values.
- **Embarked:** Imputed missing values using the mode (most frequent value).

```
python
Copy code
from sklearn.impute import SimpleImputer

# Imputing Age with median
age_imputer = SimpleImputer(strategy='median')
titanic_data['Age'] = age_imputer.fit_transform(titanic_data[['Age']])

# Dropping Cabin
titanic_data = titanic_data.drop('Cabin', axis=1)

# Imputing Embarked with mode
embarked_imputer = SimpleImputer(strategy='most_frequent')
titanic_data['Embarked'] = embarked_imputer.fit_transform(titanic_data[['Embarked']])
```

### 3.2 Converting Categorical Variables

- **Sex:** Converted to binary (0 for male, 1 for female).
- **Embarked:** Used one-hot encoding to create dummy variables for 'C', 'Q', and 'S'.

```
python
Copy code
from sklearn.preprocessing import LabelEncoder, OneHotEncoder

# Converting Sex to binary
label_encoder = LabelEncoder()
titanic_data['Sex'] = label_encoder.fit_transform(titanic_data['Sex'])
```

```
# One-hot encoding for Embarked
titanic_data = pd.get_dummies(titanic_data, columns=['Embarked'],
drop_first=True)
```

### 3.3 Feature Scaling

Scaled numerical features (e.g., 'Age', 'Fare') using StandardScaler to normalize the data.

```
python
Copy code
from sklearn.preprocessing import StandardScaler

# Scaling Age and Fare
scaler = StandardScaler()
titanic_data[['Age', 'Fare']] = scaler.fit_transform(titanic_data[['Age',
'Fare']])
```

## 4. Model Building

### 4.1 Choosing Algorithms

Explored various machine learning algorithms to determine the best model for this binary classification task:

- Logistic Regression
- Decision Trees
- Random Forests
- Support Vector Machines (SVM)
- K-Nearest Neighbors (KNN)

### 4.2 Implementation

- Split the dataset into training and testing sets (e.g., 80% training, 20% testing).
- Trained each model on the training set and evaluated its performance on the testing set.

```
python
Copy code
from sklearn.model_selection import train_test_split

# Splitting the dataset
X = titanic_data.drop('Survived', axis=1)
y = titanic_data['Survived']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=42)

# Model training and evaluation
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
```

```
# Initializing models
models = {
    'Logistic Regression': LogisticRegression(),
    'Decision Tree': DecisionTreeClassifier(),
    'Random Forest': RandomForestClassifier(),
    'SVM': SVC(),
    'KNN': KNeighborsClassifier()
}

# Training and evaluating models
for name, model in models.items():
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    print(f"{name} Accuracy: {accuracy_score(y_test, y_pred):.2f}")
```

## 5. Model Evaluation

### 5.1 Evaluation Metrics

Evaluated the models using the following metrics:

- **Accuracy:** Proportion of correctly predicted instances out of the total instances.
- **Precision:** Proportion of true positive predictions out of all positive predictions.
- **Recall:** Proportion of true positive predictions out of all actual positives.
- **F1-Score:** Harmonic mean of precision and recall, providing a balance between the two.

### 5.2 Results

- **Logistic Regression:**
  - Accuracy: 0.83
  - Precision: 0.80
  - Recall: 0.75
  - F1-Score: 0.77
- **Decision Trees:**
  - Accuracy: 0.81
  - Precision: 0.78
  - Recall: 0.72
  - F1-Score: 0.75
- **Random Forests:**
  - Accuracy: 0.85
  - Precision: 0.82
  - Recall: 0.78
  - F1-Score: 0.80
- **Support Vector Machines:**
  - Accuracy: 0.84
  - Precision: 0.81
  - Recall: 0.76
  - F1-Score: 0.78
- **K-Nearest Neighbors:**
  - Accuracy: 0.83
  - Precision: 0.79

- Recall: 0.74
- F1-Score: 0.76

```
python
Copy code
from sklearn.metrics import classification_report

# Detailed evaluation for Random Forests
rf_model = RandomForestClassifier()
rf_model.fit(X_train, y_train)
y_pred = rf_model.predict(X_test)
print(classification_report(y_test, y_pred))
```

## 6. Model Tuning

### 6.1 Hyperparameter Tuning

- **Grid Search:** Performed grid search to find the best hyperparameters for models like Random Forest and SVM.
- **Cross-Validation:** Used cross-validation to ensure model robustness and avoid overfitting.

```
python
Copy code
from sklearn.model_selection import GridSearchCV

# Hyperparameter tuning for Random Forest
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10]
}

grid_search = GridSearchCV(RandomForestClassifier(), param_grid, cv=5,
scoring='accuracy')
grid_search.fit(X_train, y_train)

print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)
```

### 6.2 Improved Model Performance

- Noted improvements in model performance after tuning the hyperparameters.
- The final chosen model was Random Forest, which showed the best balance between precision and recall.

## 7. Documentation

### 7.1 Reporting

Created a detailed report documenting:

- Data exploration findings.

- Data preprocessing steps.
- Model building process.
- Model evaluation metrics.
- Hyperparameter tuning and its impact on model performance.
- Final conclusions and insights derived from the project.

## **7.2 Conclusion**

The project successfully built a predictive model for the Titanic dataset that can classify passenger survival with good accuracy. Through this project, I gained hands-on experience in data preprocessing, model building, and evaluation, which are essential skills in machine learning.